

A Brief Tutorial on Hugging Face Transformers Lib

Hugging Face Transformers library, its associated APIs, and how to use it for fine-tuning a large language model (LLM):

Overview of the Hugging Face Transformers Library

The Hugging Face Transformers library is an open-source library that provides a large number of pre-trained transformer models and simple APIs for natural language processing (NLP) tasks. It supports various models such as BERT, GPT, T5, etc., and enables users to easily load and use pre-trained models for tasks like text classification, text generation, translation, and more. The library supports both PyTorch and TensorFlow frameworks, offering flexibility for users to choose their preferred machine learning framework.

Key APIs

- **Model Loading APIs:** APIs like `AutoModelForCausalLM`, `AutoModelForSequenceClassification`, `AutoModelForQuestionAnswering`, etc., are used to load pre-trained models for different tasks. The `AutoModel` APIs automatically select the appropriate model architecture based on the model name or configuration.
- **Tokenizer APIs:** APIs like `AutoTokenizer` are used to load tokenizers corresponding to pre-trained models. Tokenizers convert text into tokens that models can process and provide encoding and decoding functions.
- **Pipeline API:** The pipeline API abstracts the complexities of model usage, allowing users to quickly perform tasks like sentiment analysis, text generation, and named entity recognition with just a few lines of code.
- **Trainer API:** The Trainer API simplifies the training loop, handling gradient updates, evaluation, and logging. It abstracts much of the complexity behind training, making it convenient for users to fine-tune models.

Fine-Tuning a Pre-Trained Model

Environment Setup

- Install the Transformers library and its dependencies:

```
1 | !pip install transformers datasets torch
```

Load the Dataset

You can use the `datasets` library to load datasets. For example, to load the IMDb dataset for sentiment analysis:

```
1 | ```python
2 | from datasets import load_dataset
3 | dataset = load_dataset("imdb")
4 | ```
```

Load the Pre-Trained Model and Tokenizer

Choose a suitable pre-trained model and tokenizer. For instance, to fine-tune a BERT model for sentiment analysis:

```
1  ```python
2  from transformers import AutoModelForSequenceClassification, AutoTokenizer
3
4  model_name = "bert-base-uncased"
5  model = AutoModelForSequenceClassification.from_pretrained(model_name, num_labels=2)
6  tokenizer = AutoTokenizer.from_pretrained(model_name)
7  ```
```

Preprocess the Dataset

Tokenize the text data in the dataset:

```
1  ```python
2  def tokenize_function(examples):
3      return tokenizer(examples["text"], padding="max_length", truncation=True)
4
5  tokenized_datasets = dataset.map(tokenize_function, batched=True)
6  ```
```

Set Training Arguments

Define training parameters using the TrainingArguments class:

```
1  ```python
2  from transformers import TrainingArguments
3
4  training_args = TrainingArguments(
5      output_dir="./results",
6      num_train_epochs=3,
7      per_device_train_batch_size=8,
8      evaluation_strategy="epoch",
9      learning_rate=2e-5,
10     weight_decay=0.01,
11 )
12 ```
```

Initialize the Trainer

Create a Trainer instance to manage the training process:

```
1  ```python
2  from transformers import Trainer
3
4  trainer = Trainer(
5      model=model,
6      args=training_args,
7      train_dataset=tokenized_datasets["train"],
8      eval_dataset=tokenized_datasets["test"],
9  )
10 ```
```

Start Training

Call the `train` method of the `Trainer` to begin fine-tuning the model:

```
1  ```python
2  trainer.train()
3  ```
```

Evaluate the Model

After training, evaluate the model's performance on the test set:

```
1  ```python
2  results = trainer.evaluate()
3  print(results)
4  ```
```

Best Practices

- **Choose the Right Pre-Trained Model:** Select a pre-trained model that matches the task and dataset. For example, for sentiment analysis of English text, models like `bert-base-uncased` or `distilbert-base-uncased-finetuned-sst-2-english` are good choices.
- **Monitor Training Metrics:** Track metrics like loss, accuracy, precision, recall, etc., during training to assess model performance and detect issues like overfitting.
- **Adjust Hyperparameters:** Fine-tune hyperparameters such as learning rate, batch size, and number of epochs based on the task and dataset to optimize model performance.
- **Use Mixed Precision Training:** Enable mixed precision training to accelerate training and reduce memory usage.
- **Leverage Model Quantization:** Apply model quantization to reduce the model's memory footprint and improve inference speed without significantly compromising performance.

Additional Resources

- [Hugging Face Transformers Documentation](#)
- [Hugging Face Model Hub](#)
- [Pipeline Documentation](#)

The above tutorial is based on content from sources such as "*How to Use Hugging Face: Step-by-Step Guide - ML Journey*" and "*Hugging Face Transformers for NLP: A Comprehensive Guide*". For more detailed information, you can refer to the official documentation and related tutorials.

PPO in RLHF

To implement Proximal Policy Optimization (PPO) for Reinforcement Learning from Human Feedback (RLHF) using a pretrained model, follow these steps. This example assumes you have a pretrained policy model (e.g., GPT-2), a pretrained reward model, and a value network.

Step 1: Install Required Libraries

```
1 | pip install torch transformers
```

Step 2: Import Necessary Modules

```
1 | import torch
2 | import torch.nn as nn
3 | import torch.optim as optim
4 | from transformers import AutoModelForCausalLM, AutoModelForSequenceClassification,
   | AutoTokenizer
```

Step 3: Define the Policy and Value Models

```
1 | class PolicyModel(nn.Module):
2 |     def __init__(self, pretrained_model_name):
3 |         super().__init__()
4 |         self.model = AutoModelForCausalLM.from_pretrained(pretrained_model_name)
5 |
6 |     def forward(self, input_ids, attention_mask):
7 |         return self.model(input_ids=input_ids, attention_mask=attention_mask)
8 |
9 | class ValueModel(nn.Module):
10 |     def __init__(self, pretrained_model_name):
11 |         super().__init__()
12 |         self.model = AutoModelForCausalLM.from_pretrained(pretrained_model_name)
13 |         self.value_head = nn.Linear(self.model.config.n_embd, 1)
14 |
15 |     def forward(self, input_ids, attention_mask):
16 |         outputs = self.model(input_ids=input_ids, attention_mask=attention_mask)
17 |         last_hidden_states = outputs.last_hidden_state
18 |         return self.value_head(last_hidden_states)
```

Step 4: Generate Sequences with Log Probabilities

```
1 def generate_with_logprobs(policy_model, prompts, tokenizer, max_length=50):
2     inputs = tokenizer(prompts, return_tensors="pt", padding=True, truncation=True)
3     input_ids = inputs.input_ids
4     attention_mask = inputs.attention_mask
5
6     generated_tokens = []
7     logprobs = []
8     batch_size = input_ids.shape[0]
9
10    for _ in range(max_length):
11        outputs = policy_model(input_ids=input_ids, attention_mask=attention_mask)
12        logits = outputs.logits[:, -1, :]
13        log_softmax = nn.functional.log_softmax(logits, dim=-1)
14
15        # Sample next token
16        next_tokens = torch.multinomial(torch.exp(log_softmax), num_samples=1).squeeze(1)
17        logprob = log_softmax[torch.arange(batch_size), next_tokens]
18
19        generated_tokens.append(next_tokens.unsqueeze(1))
20        logprobs.append(logprob.unsqueeze(1))
21
22        # Update input_ids and attention_mask
23        input_ids = torch.cat([input_ids, next_tokens.unsqueeze(-1)], dim=1)
24        attention_mask = torch.cat([attention_mask, torch.ones(batch_size, 1)], dim=1)
25
26    generated_tokens = torch.cat(generated_tokens, dim=1)
27    logprobs = torch.cat(logprobs, dim=1)
28
29    return generated_tokens, logprobs
```

Step 5: Compute Rewards Using the Reward Model

```
1 def compute_rewards(reward_model, prompts, generated_tokens, tokenizer):
2     full_sequences = tokenizer.batch_decode(generated_tokens, skip_special_tokens=True)
3     inputs = tokenizer(full_sequences, return_tensors="pt", padding=True, truncation=True)
4     with torch.no_grad():
5         outputs = reward_model(**inputs)
6     return outputs.logits.squeeze(-1)
```

Step 6: PPO Training Loop

```
1 def ppo_train(
2     policy_model,
3     value_model,
4     reward_model,
5     tokenizer,
6     prompts,
```

```

7     epochs=10,
8     ppo_epochs=4,
9     clip_epsilon=0.2,
10    value_loss_coeff=0.5,
11    learning_rate=3e-5
12 ):
13     policy_optimizer = optim.AdamW(policy_model.parameters(), lr=learning_rate)
14     value_optimizer = optim.AdamW(value_model.parameters(), lr=learning_rate)
15
16     for epoch in range(epochs):
17         # Generate responses and log probabilities
18         generated_tokens, logprobs = generate_with_logprobs(policy_model, prompts,
tokenizer)
19
20         # Compute rewards
21         rewards = compute_rewards(reward_model, prompts, generated_tokens, tokenizer)
22
23         # Compute returns (same as rewards for simplicity)
24         returns = rewards
25
26         # Compute values for each token in the sequence
27         with torch.no_grad():
28             value_outputs = value_model(input_ids=generated_tokens,
attention_mask=torch.ones_like(generated_tokens))
29             values = value_outputs.squeeze(-1)
30
31         # Compute advantages
32         advantages = returns.unsqueeze(1) - values
33
34         # PPO updates
35         for _ in range(ppo_epochs):
36             # Recompute log probabilities with updated policy
37             with torch.no_grad():
38                 new_logprobs = generate_with_logprobs(policy_model, prompts, tokenizer)[1]
39
40             # Compute ratio and surrogate loss
41             ratio = torch.exp(new_logprobs - logprobs)
42             surrogate_loss = ratio * advantages
43             clipped_surrogate_loss = torch.clamp(ratio, 1 - clip_epsilon, 1 + clip_epsilon)
44             * advantages
45             policy_loss = -torch.min(surrogate_loss, clipped_surrogate_loss).mean()
46
47             # Compute value loss
48             current_values = value_model(input_ids=generated_tokens,
attention_mask=torch.ones_like(generated_tokens)).squeeze(-1)
49             value_loss = nn.MSELoss()(current_values, returns.unsqueeze(1))
50
51             # Total loss
52             total_loss = policy_loss + value_loss * value_loss_coeff
53
54             # Backpropagate
55             policy_optimizer.zero_grad()
56             value_optimizer.zero_grad()
57             total_loss.backward()
58             policy_optimizer.step()
59             value_optimizer.step()
60
61     print(f"Epoch {epoch}, Loss: {total_loss.item()}")

```

Step 7: Initialize Models and Tokenizer

```
1 policy_model = PolicyModel("gpt2")
2 value_model = ValueModel("gpt2")
3 reward_model = AutoModelForSequenceClassification.from_pretrained("path/to/reward/model")
4 tokenizer = AutoTokenizer.from_pretrained("gpt2")
5
6 prompts = ["Write a story about a robot who...", "Explain quantum physics in simple terms..."]
```

Step 8: Run Training

```
1 ppo_train(policy_model, value_model, reward_model, tokenizer, prompts)
```

Explanation:

1. **Policy Model:** Generates text sequences and tracks log probabilities.
2. **Value Model:** Estimates the expected reward for each token in the sequence.
3. **Reward Model:** Evaluates the quality of generated sequences.
4. **PPO Training:** Uses clipped surrogate objectives to update the policy and minimize value prediction error.

This implementation provides a foundation. For production use, handle token indexing, padding, and reward normalization carefully. Adjust hyperparameters and model architectures based on your specific use case.

SFT

Supervised fine-tuning (SFT) using Hugging Face's Transformers library and the trl library's SFTTrainer:

Step 1: Install Required Libraries

First, install the necessary libraries. If using Google Colab, you can install them directly in the notebook:

```
1 !pip install transformers datasets accelerate trl peft
```

Step 2: Load the Dataset

Choose a suitable dataset for fine-tuning. For example, to fine-tune a model for a text summarization task, you can use the XSum dataset:

```
1 from datasets import load_dataset
2
3 dataset = load_dataset("xsum")
```

Step 3: Load the Pre-trained Model and Tokenizer

Select a pre-trained model and its corresponding tokenizer. For instance, using the HuggingFaceTB/SmolLM-135M model:

```
1 from transformers import AutoTokenizer, AutoModelForCausalLM
2
3 model_name = "HuggingFaceTB/SmolLM-135M"
4 tokenizer = AutoTokenizer.from_pretrained(model_name)
5 if tokenizer.pad_token is None:
6     tokenizer.pad_token = tokenizer.eos_token
7 model = AutoModelForCausalLM.from_pretrained(model_name)
```

Step 4: Preprocess the Dataset

Tokenize the dataset and format it according to the requirements of the model and task:

```
1 def preprocess_function(examples):
2     return tokenizer(examples["document"], truncation=True, padding="max_length",
3                       max_length=512)
4
5 tokenized_dataset = dataset.map(preprocess_function, batched=True)
```

Step 5: Configure Fine-Tuning Parameters

Define the training arguments and use the SFTTrainer for fine-tuning:

```
1 from trl import SFTTrainer, SFTConfig
2
3 sft_config = SFTConfig(
4     output_dir="./smollm-sft-xsum",
5     per_device_train_batch_size=4,
6     gradient_accumulation_steps=4,
7     learning_rate=2e-5,
8     max_steps=50,
9     logging_steps=10,
10    save_steps=10,
11    fp16=True
```



```
12 )
13
14 trainer = SFTTrainer(
15     model=model,
16     args=sft_config,
17     train_dataset=tokenized_dataset["train"],
18     eval_dataset=tokenized_dataset["test"]
19 )
```

Step 6: Start Fine-Tuning

Begin the fine-tuning process:

```
1 trainer.train()
```

Step 7: Evaluate and Save the Model

Evaluate the fine-tuned model on the validation set and save the model for later use:

```
1 metrics = trainer.evaluate(tokenized_dataset["validation"])
2 print(metrics)
3 trainer.save_model("./fine_tuned_model")
```

Key Points and Best Practices

- **Choosing the Right Dataset:** The dataset should align with the specific task you want the model to perform. For example, use a dataset of customer service conversations to fine-tune a model for customer service applications.
- **Monitoring Training:** Track training metrics such as loss to ensure the model is learning effectively. Adjust hyperparameters like learning rate and batch size as needed.
- **Avoiding Overfitting:** Use techniques like early stopping and regularization to prevent overfitting. Validate the model on a held-out dataset to ensure it generalizes well.

The above tutorial is based on the official Hugging Face documentation and community tutorials. For more detailed information, you can refer to the Hugging Face LLM Course and the `trl` library documentation.

Reward Model

Reward model fine-tuning using Hugging Face. This tutorial is based on the information from Hugging Face's documentation and other relevant sources.

Step 1: Install Required Libraries

Make sure you have the necessary libraries installed. You can install them using pip:

```
1 !pip install transformers datasets accelerate trl peft
```

Step 2: Load the Dataset

The dataset for reward model training typically consists of pairs of prompts and responses, along with human preference rankings. For example, you can use the `Dahoas/rm-static` dataset.

```
1 from datasets import load_dataset
2
3 dataset = load_dataset("Dahoas/rm-static")
```

Step 3: Load the Pre-trained Model and Tokenizer

Choose a pre-trained model suitable for sequence classification. For example, using the `gpt2` model:

```
1 from transformers import AutoModelForSequenceClassification, AutoTokenizer
2
3 model_name = "gpt2"
4 model = AutoModelForSequenceClassification.from_pretrained(model_name)
5 tokenizer = AutoTokenizer.from_pretrained(model_name)
6 if tokenizer.pad_token is None:
7     tokenizer.pad_token = tokenizer.eos_token
```

Step 4: Preprocess the Dataset

Tokenize the dataset and format it according to the requirements of the model.

```
1 def preprocess_function(examples):
2     return tokenizer(examples["text"], padding="max_length", truncation=True)
3
4 tokenized_dataset = dataset.map(preprocess_function, batched=True)
```

Step 5: Prepare Training Arguments

Define the training arguments using `RewardConfig` from the `trl` library.

```

1  from trl import RewardConfig
2
3  training_args = RewardConfig(
4      output_dir="./reward_model",
5      per_device_train_batch_size=4,
6      gradient_accumulation_steps=4,
7      learning_rate=2e-5,
8      num_train_epochs=3,
9      logging_dir="./logs",
10     logging_steps=10,
11     save_strategy="epoch",
12 )

```

Step 6: Initialize the Reward Trainer

Use the RewardTrainer to fine-tune the reward model.

```

1  from trl import RewardTrainer
2
3  trainer = RewardTrainer(
4      model=model,
5      args=training_args,
6      train_dataset=tokenized_dataset["train"],
7      eval_dataset=tokenized_dataset["validation"],
8  )

```

Step 7: Start Fine-Tuning

Begin the fine-tuning process:

```

1  trainer.train()

```

Step 8: Evaluate and Save the Model

Evaluate the fine-tuned model and save it for later use.

```

1  eval_results = trainer.evaluate(tokenized_dataset["validation"])
2  print(eval_results)
3  trainer.save_model("./fine_tuned_reward_model")

```

Key Points and Best Practices

- **Data Quality:** Ensure the dataset contains high-quality human preference rankings to train an effective reward model.
- **Model Selection:** Choose a pre-trained model that is suitable for your specific task and has sufficient capacity to learn the reward function.
- **Hyperparameter Tuning:** Experiment with different hyperparameters such as learning rate, batch size, and number of epochs to optimize training results.
- **Avoid Overfitting:** Use techniques like early stopping and regularization to prevent overfitting, especially when the training dataset is small.
- **Leverage PEFT:** You can use Parameter-Efficient Fine-Tuning (PEFT) techniques like LoRA to fine-tune the reward model more efficiently. Pass a `peft_config` to the `RewardTrainer` to automatically convert the model into a PEFT model.

This tutorial provides a basic guide to fine-tuning a reward model using Hugging Face. For more detailed information, refer to the official Hugging Face documentation and related tutorials.